

Miriam Wanner

mwanner5@jhu.edu

miriamwanner.github.io

Education

- **Ph.D., Computer Science** 8/2023 - Present
M.S.E., Computer Science 8/2023 - 5/2025
Johns Hopkins University, Center for Language and Speech Processing (CLSP)
Advisors: Mark Dredze, William Walden
- **B.A., Computer Science and Mathematics** 8/2019 - 5/2023
University of Virginia
GPA: 3.96

Experience

- **Johns Hopkins University**, Baltimore, MD 8/2023 - Present
Graduate Student Researcher, Advisors: Mark Dredze, William Walden
- **Microsoft Research**, Redmond, WA 6/2025 - 8/2025
Research Intern, Advisor: Nick Craswell
- **Castorani Lab at the University of Virginia**, Charlottesville, VA 1/2021 - 5/2023
Undergraduate Student Researcher, Advisor: Max Castorani
- **REU at George Mason University**, Fairfax, VA 6/2021 - 8/2021
Undergraduate Student Researcher, Advisor: Antonios Anastasopoulos

Publications and Preprints

1. **Miriam Wanner**, Mark Dredze, William Walden. On the Threat Model of Weird Generalization and Emergent Misalignment. *Preprint*, 2026
2. **Miriam Wanner**, Hannah Collison, William Jurayj, Benjamin Van Durme, Mark Dredze, William Walden. Weird Generalization is Weirdly Brittle. *Preprint*, 2026.
3. William Walden & **Miriam Wanner**. Reasoning Models Will Sometimes Lie About Their Reasoning. *Preprint*, 2026.
4. Sathvik Napa, Utkarsh Singh, Chengyuan Xue, **Miriam Wanner**, William Walden. Can Reasoning Models Detect Changes to their Chains of Thought?. *EMNLP*, 2026
5. **Miriam Wanner**, Leif Azzopardi, Paul Thomas, Soham Dan, Benjamin Van Durme, Nick Craswell. All Claims Are Equal, but Some Claims Are More Equal Than Others: Importance-Sensitive Factuality Evaluation of LLM Generations. *EMNLP Findings*, 2026.
6. William Walden, Kathryn Ricci, **Miriam Wanner**, Zhengping Jiang, Chandler May, Rongkun Zhou, Benjamin Van Durme. How Grounded is Wikipedia? A Study on Structured Evidential Support. *ACL Findings*, 2026.
7. **Miriam Wanner***, Sophia Hager*, Anjalie Field. Does Local News Stay Local?: Online Content Shifts in Sinclair-Acquired Stations. *NLP4CSS*, 2026.

8. **Miriam Wanner**, Benjamin Van Durme, Mark Dredze. DnDScore: Decontextualization and Decomposition for Factuality Verification in Long-Form Text Generation. *EMNLP*, 2025.
9. Jiefu Ou, William Gantt Walden, Kate Sanders, Zhengping Jiang, Kaiser Sun, Jeffrey Cheng, William Jurayj, **Miriam Wanner**, Shaobo Liang, Candice Morgan, Seunghoon Han, Weiqi Wang, Chandler May, Hannah Recknor, Daniel Khashabi, Benjamin Van Durme. ClaimCheck: How Grounded are LLM Critiques of Scientific Papers? *EMNLP Findings*, 2025.
10. Zhengping Jiang, Jingyu Zhang, Nathaniel Weir, Seth Ebner, **Miriam Wanner**, Kate Sanders, Daniel Khashabi, Anqi Liu, Benjamin Van Durme. Core: Robust Factual Precision Scoring with Informative Sub-Claim Identification. *ACL Findings*, 2025.
11. **Miriam Wanner***, Seth Ebner*, Zhengping Jiang, Mark Dredze, Benjamin Van Durme. A Closer Look at Claim Decomposition. **SEM*, 2024.
12. **Miriam S. Wanner**, Jonathan A. Walter, Daniel C. Reuman, Tom W. Bell, & Max C.N. Castorani. Dispersal synchronizes giant kelp forests. *Ecology*, 2024.
13. Nathaniel Krasner*, **Miriam Wanner***, Antonios Anastasopoulos. Revisiting the Effects of Leakage on Dependency Parsing. *ACL Findings*, 2022.

Awards and Honors

- Graduation With Distinction, University of Virginia 5/2023
- College Science Scholar, University of Virginia 9/2019 – 5/2023
- Echols Scholar, University of Virginia 9/2019 – 5/2023

Teaching Experience

- Johns Hopkins University 1/2025 - 5/2025
Teaching Assistant, *NLP for Computational Social Science*
- University of Virginia 1/2022 - 12/2022
Teaching Assistant, *Introduction to Programming*
- University of Virginia 8/2020 - 12/2022
Teaching Assistant, *Discrete Mathematics*

Skills

- Programming languages: Python, C++, Bash, R
- Frameworks/Toolkits: Claude Code, PyTorch, HuggingFace Transformers, $\text{\LaTeX}$, Git, Slurm
- Coursework: Machine Learning (ML), Deep Learning, Natural Language Processing (NLP), NLP for Computational Social Science, NLP: Self-Supervised Models, Human-in-the-Loop ML, Theory of Replicable ML, Machine Social Intelligence

Service

- Reviewer: ARR 2025 (secondary), ICWSM 2026, ARR 2026, BlackboxNLP 2026
- JHU CLSP Committees: Social (2023-24), Recruitment (2024-25), Graduate Admissions (2025-26)
- Application Mentor, JHU CLSP pre-application support program (2024, 2025)